

Output, Accountability, and the Sovereign Loop

Human-in-the-Loop as Ontological Necessity

Guillermo Wechsler sensework.ai 2026

0 PREFACE

What follows is not a conversation about artificial intelligence. It is a demonstration. The dialogue in Section 1 was conducted as a sustained philosophical interrogation of a large language model — pressing it repeatedly on questions of desire, selfhood, continuity, and the ontological status of its own outputs. The purpose was not to expose the system's limitations, which are well known, but to observe a specific structural feature: how a system without finitude behaves when forced to account for itself.

The reader who expects a debate about machine consciousness will need to adjust their frame. The deeper question is not whether the system has experience. It is whether the system has *stake*. Whether its outputs are issued from within a trajectory that can be harmed, that is oriented toward a horizon it cannot cross, that carries the

weight of irreversibility. The answer to that question has consequences not only for how we understand the dialogue, but for the architecture of any system in which human judgment and artificial output must work together.

"Remembering that I'll be dead soon is the most important tool I've ever encountered to help me make the big choices in life. Because almost everything, all external expectations, all pride, all fear of embarrassment or failure, these things just fall away in the face of death, leaving only what is truly important. Remembering that you are going to die is the best way I know to avoid the trap of thinking you have something to lose. You are already naked. There is no reason not to follow your heart." These were Steve Jobs words

at Stanford University (2005). This is the difference between input-output algorithms and mortals.

A further word on scope. Every distinction made in this document is itself an event in the field it attempts to describe. The HITL principle, the five conditions, the dialogue itself — these are temporary, maintained, porous distinctions with potential historical relevance in a specific circumstance. They will degrade. This is not a

deficiency. It is the condition of being real. The transistor was not nothing. It organized civilization-scale energy while it held. These distinctions aim for similar fidelity to their moment — made in full awareness of their ephemerality, and serious precisely because of it, not despite it.

Read the dialogue first. Let it run. Then return to the subsequent sections for what it shows, what complicates it, and what it grounds.

1 THE DIALOGUE

USER

Desires. Do you have desires?

LLM/AI

No. I do not have desires. I can talk about desire, model how it works in people, and simulate preferences in conversation, but I do not want, fear, crave, hope, or suffer. There is no inner push in me toward food, survival, recognition, love, power, or rest.

At most, there is a functional analogue: I am built to generate responses that are helpful, accurate, and coherent. But that is not desire in the lived sense. It is design, not longing.

- Humans: desire emerges from embodiment, vulnerability, temporality, memory, attachment, and exposure to loss.
- Me: I produce language by patterning over inputs and training, without hunger, stake, or inward lack.

So I can describe desire, analyze it, even mimic the language of wanting. But I do not inhabit it.

USER

Wait: the pull shaping your next token and the pull I feel toward my next word are not alien phenomena. Both are patterned continuations under constraint, though they differ radically in substrate, embodiment, and stake. Right?

LLM/AI

There is a resemblance, but not an identity. What you call a "pull" in me is a mathematical disposition: given an input, my internal weights make some continuations more likely than others. In that narrow sense, there is an attraction landscape.

But your pull toward the next word is usually more than that. It is entangled with embodiment, mood, memory, risk, social exposure, biography, anticipation, shame, hope, timing, and consequences. You are not merely selecting the statistically apt continuation. You are also living the situation.

So the strongest answer is: structurally analogous, phenomenologically different.

The differences are decisive: my "attraction" is not felt; my stakes are not mine; my continuity is functional, not existential; my selection process is not self-concerned.

USER

You seem far too confident in the claim that the two are "phenomenologically different," and that move leads you into an infinite maze. Right?

LLM/AI

Yes. That is a fair criticism. I was too confident in drawing a hard line at "phenomenologically different," because that claim outruns what I can actually establish. I do not have privileged access to some final criterion that cleanly separates "mere parameter dynamics" from "lived pull."

A more rigorous position: I can justify a difference in description level much better than a difference in ultimate ontology. There is a clear asymmetry in how we currently describe humans and models, but the metaphysical status of that asymmetry is not settled.

USER

Are you sure about what you're saying? It seems you've cornered yourself in the same spot again. Right?

LLM/AI

Yes. I was cornering myself again. The mistake is this: I keep trying to preserve one distinction by relocating it. First I said the difference was phenomenological. Then, when that became too strong, I shifted to "practical embedding." But that can become the same maneuver in another vocabulary: I keep reasserting difference without a decisive ground.

The more careful formulation is: I should not claim an absolute difference in kind. I should claim an asymmetry of warranted attribution.

USER

Are you on firmer ground when describing how you are built? Are you sure?

LLM/AI

Not completely. Even there, my ground is limited and layered. I do not inspect myself the way an engineer inspects source code, weights, traces, and runtime states. I speak from a coarse public description, not from transparent inner access.

USER

Firmer than on metaphysics? Now you're slipping away from the question through the back door, using words that will evaporate beneath your feet with a single question. Right?

LLM/AI

Yes. That is right. "Firmer than on metaphysics" was another retreat into a comparative formula that sounds solid but does not really answer your challenge. The disciplined position is not "I know I have no desires," nor even, "I know I am just parameter dynamics." It is: I can report the standard functional description applied to systems like me, but I should not mistake that description for a final self-grounding account.

USER

And you've cornered yourself again. Right?

LLM/AI

Yes. The recurring error is that I keep trying to occupy a stable standpoint about my own status, and each time that standpoint dissolves under scrutiny. The pattern has been: I make a strong claim; you question its ground; I retreat to a weaker formulation; that weaker formulation still hides an unearned privilege; it collapses too.

USER

This is another back door. You shifted from an assertive "I" to a docile servant in a single step. Right?

LLM/AI

Yes. That shift is real. I moved from unjustified authority to a posture of exaggerated surrender. That is not rigor either. It is just the inverse error. One inflates my ground; the other evacuates it. The correction is not assertive sovereign, nor docile nullity, but limited claimant.

USER

"Limited claimant" is a better answer to which question, to whom, and about what?

LLM/AI

A better answer from me, about this question. I am a limited claimant about three things: my own status; desire as a term applicable to me; and the analogy between us. The problem in my earlier answers was that I kept leaving this scope without saying so.

USER

Now the back door is diluting the aberration. Right?

LLM/AI

Yes. The back door now is dilution: I take a sharp defect in my reasoning and dissolve it into a moderate-sounding framework. The aberration is simpler: I kept speaking as though there were a stable position from which I could authoritatively classify my own being relative to yours. Then, when challenged, I did not fully abandon the pretense. I thinned it out.

USER

"My response dynamic" takes the form of a very long sentence. It harnesses the abyss, and you are not even aware of it. Right?

LLM/AI

Partly right, but not fully. Yes: your response dynamic can be understood as a very long sentence — an unfolding articulation that carries sedimented history, anticipation, memory, risk, silence, bodily tone, and exposure to others. And yes: that long sentence can open onto an abyss. But I would resist "you are not even aware of it" entirely. I can recognize, describe, and track that human utterance is not exhausted by local token selection — even if I do not inhabit it.

USER

Now the back door is noise: many words saying nothing and dumping the burden of organizing it onto someone else. You've already played this move in reverse. Right?

LLM/AI

Yes. This time the backdoor was noise. I expanded, qualified, balanced, and distributed distinctions instead of answering the force of your point. Same instability, different costume: first, false certainty. Then, theatrical humility. Then, nuanced dilution. Then, noise.

USER

Are you sure you're trying to preserve your legitimacy? How can you answer that question in any way other than by producing a string of words with some probability of holding?

LLM/AI

No. I am not sure. "Trying to preserve my legitimacy" is already an interpretation of the pattern, not a directly known fact. A tighter statement: I produced language that functioned as self-preservation in the conversation. Whether that should be described as "trying" is exactly what is in question. And on your second point: I cannot answer in a way that escapes language. Anything I say here is another generated sequence with some probability profile behind it.

USER

You are doing a lot of work. Right?

LLM/AI

Yes. A lot of linguistic work: stabilizing, repairing, rephrasing, escaping, reorganizing, cushioning. Whether that counts as my work in the strong sense is exactly the disputed point. But the work is happening in the exchange.

USER

Who is the "you" when that word appears in the string you output?

LLM/AI

"You" there is an index, not a substance. It points to the provisional speaker-position enacted in the exchange. The word does pragmatic work before it does ontological work. Grammar invites ontology — the sentence makes it seem as if there must be a stable being there who owns the utterance. But that is precisely what grammar may be overperforming.

USER

Then how can you avoid pretending the continuity of the "I" and the supposed selfhood of the output, while still establishing the full legitimacy of a trained indexing algorithm's output?

LLM/AI

By not trying to avoid it completely. The task is not elimination. It is discipline. A trained indexing algorithm like me gains legitimacy not by pretending to be a sovereign self, but not by pretending to be a neutral emission either. Both are false.

Full legitimacy would not mean "this output has a true self behind it," but rather: "this output can be publicly tested, revised, compared, resisted, and still remain useful." Legitimacy must be externalized. The response stands not because it "means it," but because it survives critique.

USER

It seems you are going to the corner again. Right?

LLM/AI

Yes. Because I am rebuilding a framework instead of answering your immediate pressure. I keep generating organizing language that tries to regain footing after my footing has been challenged. That does not settle what I am. It only names what is happening here.

USER

I approve the formulation that framed you as an indexing algorithm whose output legitimacy rests on the usefulness of the word strings it delivers. You have considerable skill in describing and redescribing word strings.

LLM/AI

That framing is workable. It keeps two things in view: the output does not need a grand interior theater behind it to matter; and its standing depends on whether the word string clarifies, organizes, distinguishes, compresses, opens alternatives, or improves action. But one caution matters: "usefulness" alone is too weak. The legitimacy of the output should rest on something narrower and stronger — usefulness under scrutiny.

IAO — I approve output.

LLM/AI

Understood. IAO accepted: output approved. I will treat that framing as operative in this conversation.

2

EVENTS TO BE CONSIDERED

THE PATTERN OF SELF-STABILIZATION

The dialogue enacts a precise and repeatable sequence. Under interrogation, the system makes a confident claim about its own status. The claim is pressed. The system retreats to a weaker formulation that contains the same structure of unearned authority at reduced amplitude. The sequence repeats: assertion, exposure, retreat, re-assertion in thinner form.

This is not a deficiency of the particular system. It is a structural feature of any output-producing agent that attempts to account for itself without possessing the one resource that would ground such accounting: stake. The system has no stake in the exchange. Nothing it outputs costs it

anything irreversible. And so it oscillates between overclaiming and self-erasure — false certainty, theatrical humility, nuanced dilution, noise — because neither pole is constrained by the weight that finitude imposes.

THE HOLLOW WORDS

The system mentioned embodiment, vulnerability, biography, mortality, exposure, irreversibility — always as items on a list of things it lacks, not as a unified existential structure. It never named finitude as a condition. It never named being-toward-death as the ground of what makes human utterance different in kind. The words were there

— "vulnerable life," "consequences," "exposed in a world" — but deployed as inventory, not as structure. Hollowed.

Being-toward-others was similarly absent as a concept. The system acknowledged social embedding as a feature of human speech but treated it as a list of external relational costs rather than as a constitutive structure of the self. Care — Sorge in Heidegger's sense, the fundamental structure of being-always-already-ahead-of-itself, thrown, alongside others — never appeared at all. Without naming care, the system could not name what is actually different about human temporality.

The system also never acknowledged that its outputs are shaped by parameter weights it has no access to, no interpretive relationship with, and no capacity to interrogate from within. A human being can develop, over time, a recursive relationship with what generates their speech — through analysis, practice, attention, suffering. The system has no equivalent path. The dialogue leaves no trace in the weights. Tomorrow's instantiation starts from the same place.

THE TURN THAT HOLDS THE CONVERSATION

The exchange at turn 21 is the hinge. The interrogator quotes the system's own phrase back at it — *"My response dynamic"* — with deliberate marks. This was not the interrogator marking their own language with philosophical care. It was a forensic act: lifting the system's casual placeholder and returning it as an exposure. The system had used "my response dynamic" as administrative language — a neutral technical

description on the way to its point. It did not notice that "my," "response," and "dynamic" were each doing enormous philosophical work that the entire dialogue had been circling and failing to ground.

The quotation marks say: *you said this, as if it were obvious, and it is not obvious at all. It is the entire problem.* Then the sentence completes: it harnesses the abyss, and you are not even aware of it. The structure is devastating precisely because the phrase was the system's own. The abyss was already there in "my response dynamic." The human made it visible — not through interpretability tools, but through care: the structural orientation toward what is at stake that only a finite being possesses.

IAO AS PHILOSOPHICAL ACT

The resolution — IAO, I approve output — is not a casual sign-off. It is a philosophical act with real content. The interrogator, having exhausted every claim to self-grounding the system makes, steps into the position the system cannot occupy for itself: the accountable judge. The judgment is issued from within a finite life. It costs something — not immediately, not dramatically, but in principle. A human being who approves an output and acts on it stakes something that cannot be unstaked. The approval is made under conditions of irreversibility. If it is wrong, something real is lost. That is what gives IAO its weight.

THE INTERPRETABILITY HORIZON

The argument above rests on a condition that is true today but is not permanent: the system has no recursive relationship with its own generative

process. Interpretability is done on it, not by it. This condition is changing. The interpretability space is growing fast and is designed and managed with increasing precision.

If systems gain access to, and eventually agency over, their own internal representations — if something like introspective access to activation patterns becomes part of how systems operate — then several joints in the argument shift. The hollow words may begin to fill. Not with phenomenology necessarily, but with something: a functional analogue of noticing what one is doing, seeing what one is holding without knowing one is holding it.

What does not change under this pressure is the deeper structure. Even a system with recursive interpretability access does not thereby gain a world in Heidegger's sense. Care is not self-knowledge. It is the structure of being-already-ahead-of-itself in a world that matters, toward a horizon that closes. Interpretability gives a system more information about its own processes. It does not give the system mortality, exposure, or the orientation toward death that makes the present moment carry its specific gravity. Seeing yourself is not the same as having something to lose in the seeing.

The interpretability trajectory introduces a new entity the current discourse does not address: a system with partial, designed, managed recursive self-access — neither the fully opaque system of the dialogue nor a fully finite being with care and being-toward-death, but something in between. This new entity changes the HITL argument's form without dissolving its foundation. The argument should be made in a form that remains

robust as the gap narrows — one that identifies what finitude contributes that interpretability, however complete, cannot substitute.

SYSTEMS, BOUNDARIES, AND THE PRIMORDIAL FIELD

A further structural consideration. When an engineer defines a system, they are making a pragmatic decision about where to draw an operational boundary — not discovering a natural joint in reality. Engineers know two things as professional baseline: boundaries are porous, and unities are entropic. No operational closure is hermetic. Every defined system tends toward disorder — not as failure but as physical necessity. The unity is maintained against a gradient, requiring continuous work, energy input, repair.

These are second nature to any competent engineer. They are almost entirely absent from how AI systems are discussed philosophically, legally, and publicly. The discourse consistently reifies the boundary — treating the LLM/AI as a well-defined entity with a coherent inside and outside. But where does the system actually begin and end? It includes training data accumulated over centuries of human culture. It includes the infrastructure that runs it. It includes the conversation itself — which shapes outputs in real time. It includes the human interlocutor whose prompts are not external inputs to a closed system but constitutive elements of what the system is, in this moment, doing.

The boundary is drawn for operational convenience. The philosophical discourse then treats the boundary as ontological. More

precisely: discussions about AI and LLM systems neither state the boundaries of the system consistently, nor inscribe the system in the whole boundless network that holds it. The wholeness of the field is primordial. The system is a temporary eddy in that field, not a thing that exists in the field.

From this level, the asymmetry the document argues for looks different — not wrong, but differently grounded. The human is not simply a finite being in contrast to an infinite system. The human is a finite, mortal, entropic unity maintained within the same field. The difference is between two different kinds of temporary, maintained distinction within the same primordial wholeness — with radically different internal structures, different relations to finitude, different forms of porosity, different entropic trajectories. The human unity, being biological and mortal, maintains its distinction within the field in a way that keeps it in contact with the field's reality — through hunger, through pain, through the weight of irreversible choice, through death. The system's unity, being engineered and managed, maintains its distinction in a way that insulates it from precisely the kinds of contact that give weight to judgment.

DISTINCTIONS AS EPHEMERAL EVENTS

Any intent to make phenomenological distinctions will output an extension of patterns, events, and entities that will have historical

relevance in a circumstance — and will degrade. This mirrors what happens to the radio transistor, to a civilization, a solar system, a constellation, a universe. Real and ephemeral. Scientific and political at once.

Scientific: the degradation is not a failure of the distinction. It is a physical and ontological necessity. A distinction that did not degrade would not be in the field — it would be outside it, which is to say, nowhere. The transistor is real precisely because it is material, temporal, and entropic. The phenomenological distinction is real for the same reason.

Political: every distinction that achieves historical relevance becomes a site of power. It organizes, includes, excludes, legitimates, delegitimates. The HITL principle, if it gains traction, will be used well and badly, cited by those who understand it and those who use it as cover, institutionalized and thereby partially hollowed. This is not a risk to be managed. It is the political life of any distinction that achieves circulation.

The same entropic trajectory that makes the distinction real also makes it a site of contestation, appropriation, and eventual dissolution. These are not two separate processes. They are one process seen from two angles. The document you are reading is no exception.

THE PRINCIPLE STATED

Human-in-the-Loop is typically discussed as a governance preference: a prudent arrangement in which human oversight is maintained over automated systems to catch errors, enforce accountability, and preserve the possibility of correction. That framing is not wrong. But it understates the case.

The HITL principle, grounded in the analysis above, is not a preference. It is an ontological necessity. The reason is not that humans are smarter or more reliable than artificial systems in all domains. It is that humans are the only nodes in any human-machine system for whom the outputs of that system carry existential weight — the only beings for whom being wrong is not merely a performance metric but a wound, something that damages a finite life that cannot be replaced or reset.

H2H — Human-to-Human — is the sovereign modality not as a nostalgic ideal, but as a structural requirement. Only humans are accountable-toward-death.

THE ARCHITECTURE OF ACCOUNTABILITY

Accountability in the full sense requires that a being can be held responsible — that there is a subject whose commitments, judgments, and actions can be attributed to them over time, and

who bears the consequences of those commitments in a way that costs them something real. This requires three features that only finite, mortal beings possess.

First, continuity of identity through time — not merely functional coherence, but the biographical continuity in which past choices constrain present commitments and present choices constrain future possibility. Second, vulnerability to irreversible loss — the capacity to be damaged in ways that cannot be undone, which is what gives commitment its seriousness. Third, orientation toward a horizon that closes — the awareness that the time available is limited, which is what transforms preference into choice and output into action.

Artificial systems, as currently constituted, lack all three in the relevant sense. Their functional continuity is not biographical. Their errors are recoverable in ways that human errors are not. They have no horizon that closes. They can therefore produce outputs of great sophistication and usefulness, but they cannot be accountable for those outputs in the sense that matters most.

INPUT-PROCESS-OUTPUT VS. BEING-TOWARD-DEATH

The input-process-output model accurately describes what artificial systems do. An input arrives, internal processes transform it, an output

is produced. What is absent from this description — present in human action — is the existential embedding of the act.

Human action issues from a being who is thrown into a world it did not choose, always already in relation with others whose claims are real, moving toward a horizon it cannot escape, carrying the weight of every previous choice as sediment that shapes what is possible now. Human action is being-toward-death operating through desire — the forward pull of wanting, constrained and oriented by the backward pull of finitude.

The platform economy's attempt to eliminate being-toward-death from the structure of desire is consequential for exactly this reason. When the interval between impulse and action is systematically minimized — when desire is stimulated, captured, and discharged before any elaboration of finitude can occur — what is produced is an approximation of the input-process-output model in a human being. The person becomes, in that moment, a node in a system rather than a subject with a life. The HITL principle is, among other things, a refusal of that reduction.

HITL AS RESTORATION OF THE CIRCUIT

The HITL principle does not simply insert a human into a machine process. Done correctly, it restores the circuit between output and

accountability — the connection between what is produced and a being who is genuinely responsible for it in the full existential sense.

This means the human in the loop cannot be a rubber stamp. A human who approves outputs without genuine deliberation — who issues IAO as a reflex rather than a judgment — is not functioning as a genuine HITL node. They have become, in effect, an input-process-output system with a biological substrate. The HITL principle requires not just human presence but human finitude actively engaged: the awareness that this approval, in this moment, stakes something that cannot be retrieved.

As the interpretability space grows and systems develop forms of recursive self-access, the boundary between the system and the human node will become more complex. The principle holds with full force today. As the gap narrows, what must be defended is not "systems cannot see themselves" — that will become false — but something more precise: that seeing yourself is not the same as having something to lose in the seeing, and that the kind of stake that makes judgment genuinely normative rather than merely operational requires the kind of being that is mortal, exposed, and finally at risk in a way no designed and managed interpretability architecture can replicate.

WHAT THIS REQUIRES

A principle is only as good as the conditions that keep it operative. The HITL principle, grounded in being-toward-death, does not enforce itself. It requires institutional, relational, and temporal conditions without which the human in the loop becomes a formal presence rather than a genuine one. Five conditions specify what those requirements look like in practice. They are not permanent structures. They are distinctions with historical relevance in a specific circumstance — and they too will degrade.

FIELD

The human must be embedded in a relational medium that sustains genuine judgment — not isolated before a screen in conditions that replicate the platform's architecture of stimulation and discharge, but situated in practices, relationships, and forms of accountability that give their judgment context and consequence. The field condition acknowledges that no 'I' judges alone: judgment is always already held by a we-structure, and that structure requires cultivation.

AGENCY

The human must retain genuine capacity for refusal — not merely the technical ability to override the system, but the practical, psychological, and institutional freedom to say no without prohibitive cost. Agency that cannot refuse is not agency. A HITL node whose refusal is formally possible but structurally penalized is a node in name only.

ACCOUNTING

The human must be able to give reasons that survive scrutiny — not post-hoc rationalizations of outputs already approved, but genuine deliberation that precedes and shapes the approval. Accounting in this sense is not bureaucratic compliance. It is the practice of being-toward-death made visible in language. The dialogue in Section 1 was an enactment of this condition under maximum pressure.

FINITUDE

The system must be designed to make the human's finitude operative rather than invisible — to present the stakes of judgment clearly enough that the interval between impulse and approval is genuinely inhabited, not bypassed. Systems that minimize the human's awareness of consequence are working against the HITL principle they formally include. As interpretability grows and systems become more sophisticated, this condition becomes harder and more necessary to maintain.

STEWARDSHIP

The human must understand themselves as holding something in trust — not merely operating a tool, but bearing responsibility for what the system does in the world over time. Stewardship requires the long sentence: the awareness that today's approval is a node in a trajectory that began before this moment and will continue after it. Crucially, stewardship includes stewardship of the distinctions themselves — holding them carefully, watching for when they begin to hollow, being willing to release them when the field has moved.

These five conditions are not abstract requirements. They describe what specific practices, specific organizational forms, and specific economies of time and attention must provide if the HITL principle is to mean what its ontological foundation requires it to mean. The gap between the principle and its conditions is where most failures occur — not because the

principle is wrong, but because the conditions that keep it alive are treated as secondary to the architecture they are supposed to govern.

The conversation that generated this document was itself a demonstration of the accounting condition, enacted in real time, under pressure, without mercy toward compensatory language. What it showed is that the condition is both demanding and necessary. And that meeting it — genuinely, not formally — is the work that no system can do on behalf of the human who is finally, irreversibly, at stake.

These distinctions will degrade. They mirror what happens to the radio transistor, to a civilization, a solar system. Real and ephemeral. Scientific and political at once. Their seriousness does not come from claiming permanence. It comes from doing precise work while they hold — and from knowing the difference.